

CS 214 Recitation(Sec. 6)

Zuohui Fu

Ph.D. Department of Computer Science

Office hour: Mon 2pm-3pm

Email: zf87 AT cs dot rutgers dot edu

11/02/2017

Topics

- Solution of HW 5
- Thread and Process
- Thread creation and joining

Solution-5-process information

- How to get process information
 - The information for each process is stored in a folder name by its PID.

```
[~sh-4.2$ ls /proc
1      118      16054     18785     26531     29152     4181      5444      7700      9732
10     11824     161       18857     26536     29153     4185      5447      771       9734
100    11836     16211     18858     26537     2922      4188      55        772       98
10047  11841     16213     18880     26539     29229     4189      5519      773       9810
10060  11845     1623      1893      26541     29403     419       552       774       9814
10126  11848     16237     1895      26549     29571     4192      5523      775       9819
10132  11849     16336     19        26550     29687     420       5525      7785      9874
10136  11850     16345     19101     26552     29704     421       5529      78        99
10140  11851     16392     19209     26560     29870     422       5565      783       9952
10152  119      165       19219     26596     3        4251      5567      784       ACPI
10162  11978     16537     19220     26601     30200     4259      5581      785       asound
10164  12       16618     19225     26606     30213     4262      5585      79        buddyinfo
1017   120      167       19523     26608     30280     4279      5590      7948      bus
10183  12121     16713     1974      26610     30285     4288      5597      7953      cgroups
10195  12192     16717     2        26623     30295     43        5609      7959      cmdline
10215  123      1683      20165     26629     30564     4328      5614      7964      consoles
```

Solution-5-process information

For more, please visit:

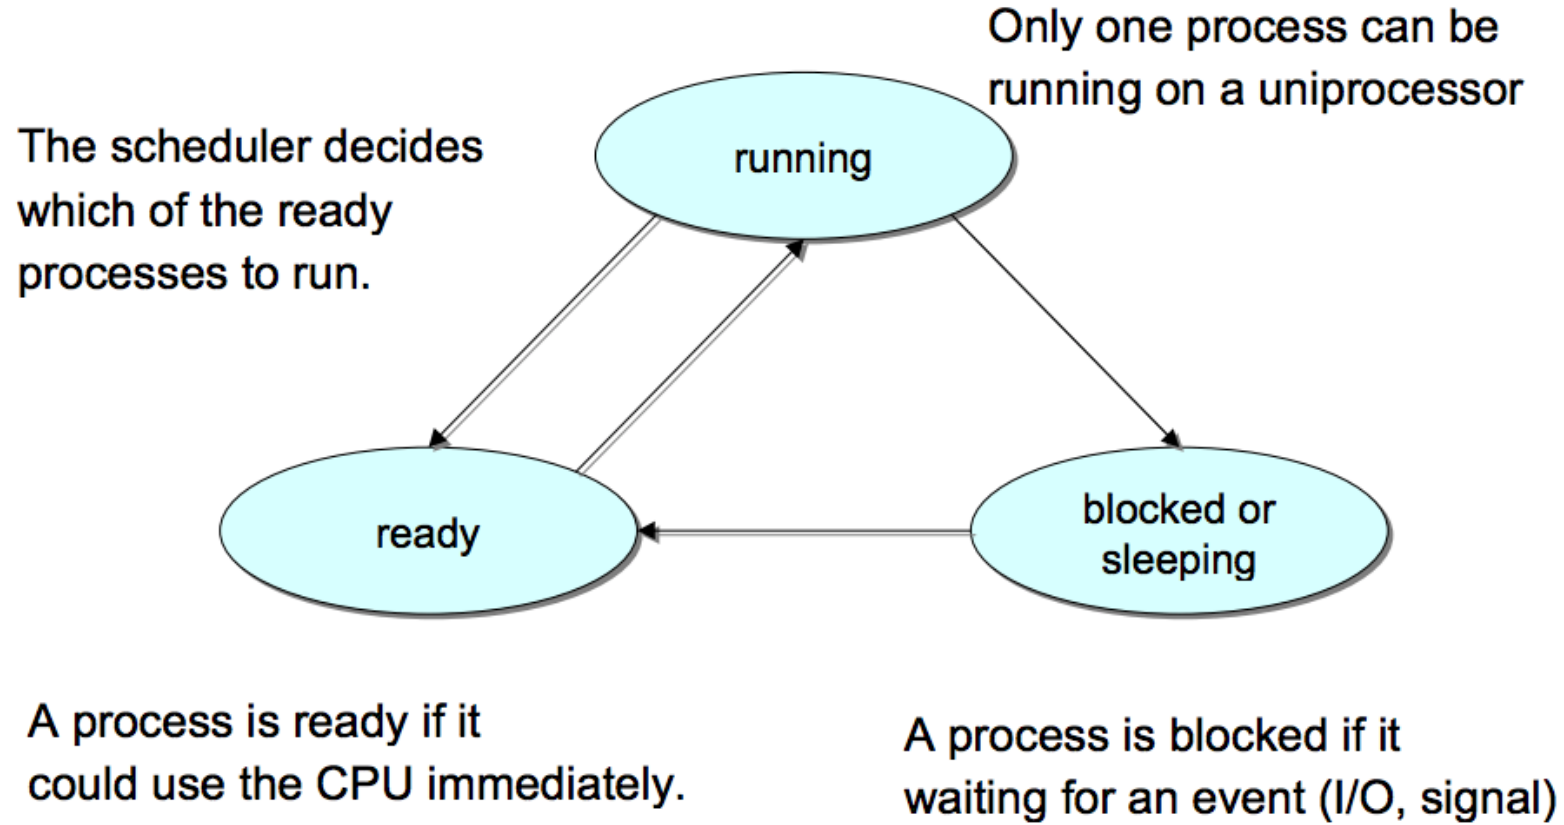
<https://www.linux.com/news/discover-possibilities-proc-directory>

<http://man7.org/linux/man-pages/man5/proc.5.html>

Ex: file **'status'** (comm, schedstat)

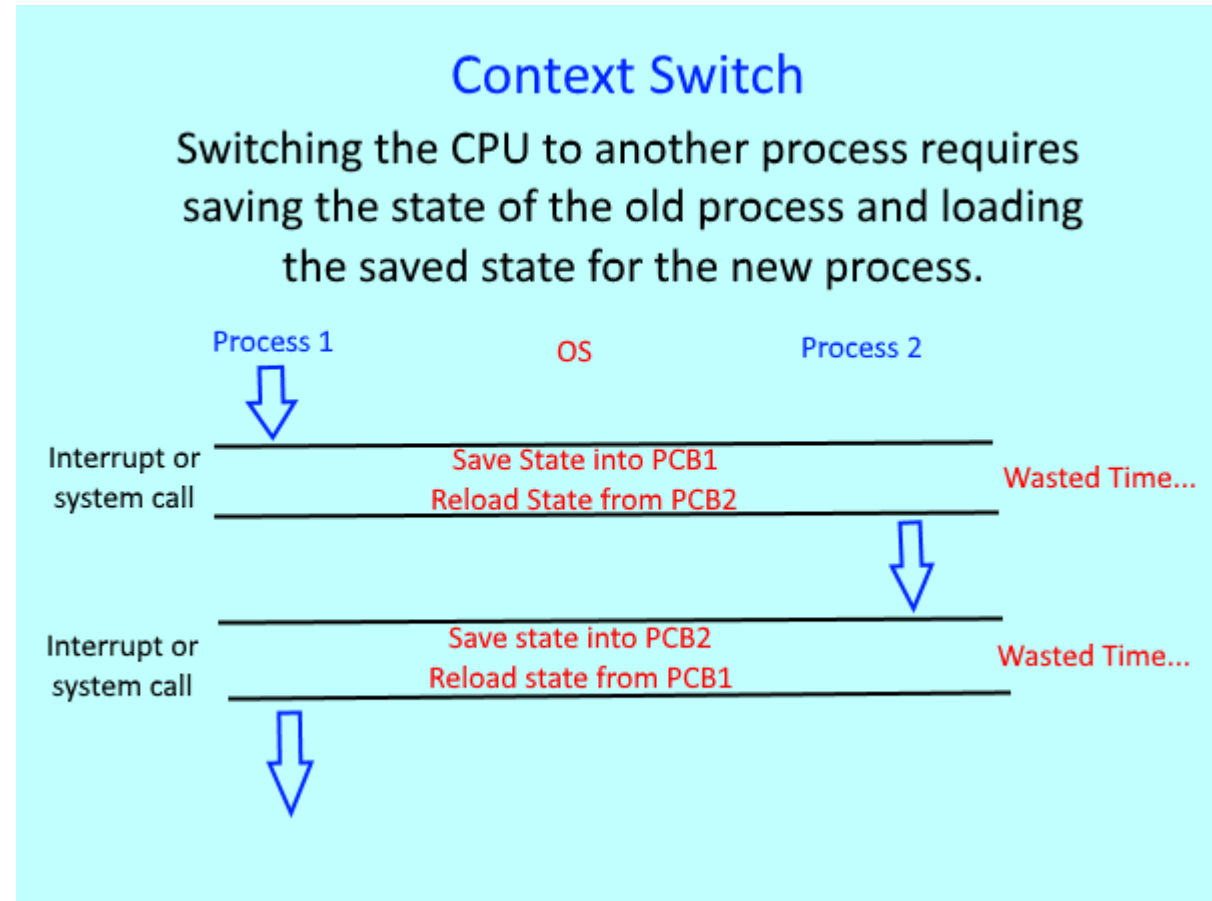
```
[~sh-4.2$  
[-sh-4.2$  
[-sh-4.2$  
[-sh-4.2$ cat /proc/187/status  
Name:   deferwq  
State:  S (sleeping)  
Tgid:   187  
Ngid:   0  
Pid:    187  
PPid:   2  
TracerPid: 0  
Uid:    0      0      0      0  
Gid:    0      0      0      0  
FDSize: 64  
Groups:  
Threads: 1
```

Process

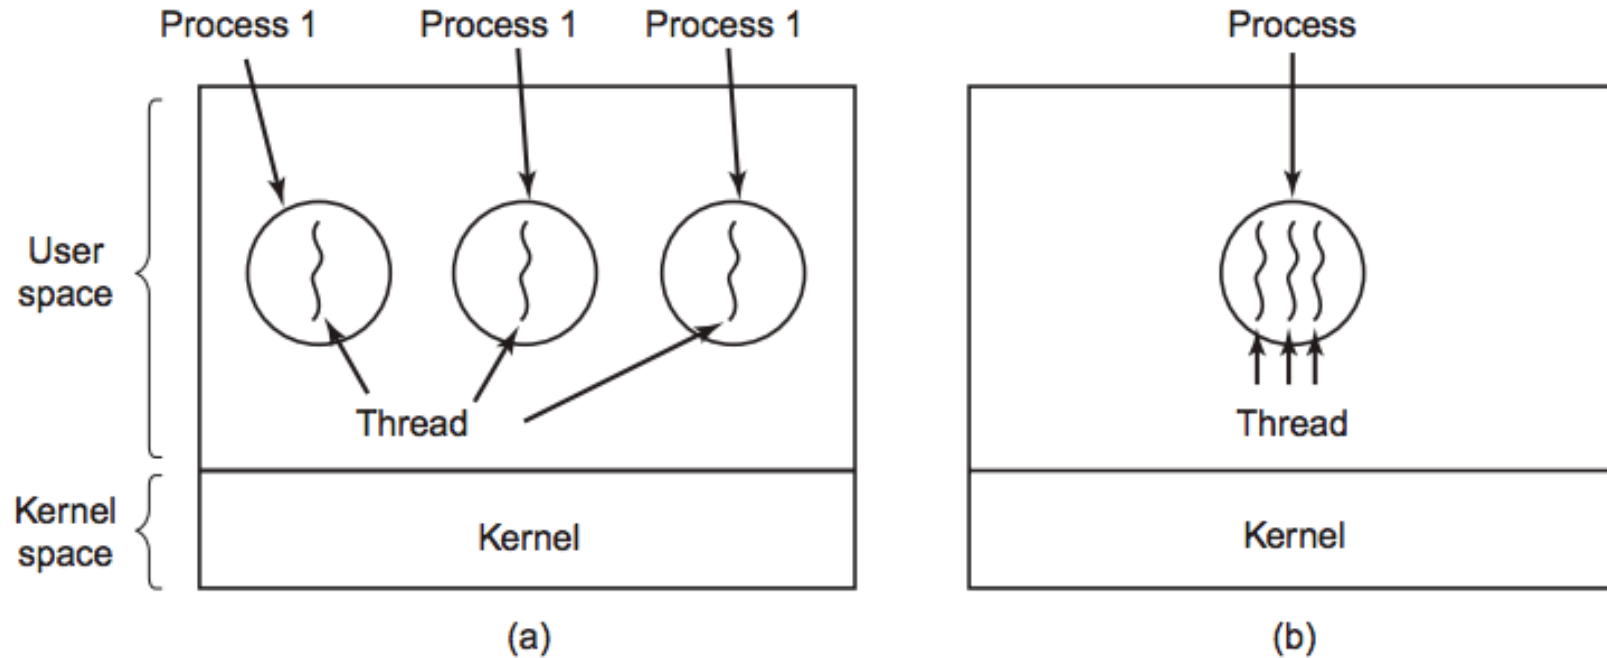


Context Switch

A context switch is the computing process of storing and restoring the state (context) of a CPU such that multiple processes can share a single CPU resource. The context switch is an essential feature of a multitasking operating system.



Properties of threads



(a) would be used when the three processes are essentially unrelated, whereas (b) would be appropriate when the three threads are actually part of the same job and are actively and closely cooperating with each other.

PCB

PCBs are data structures:

- dynamically allocated inside OS memory

When a process is created:

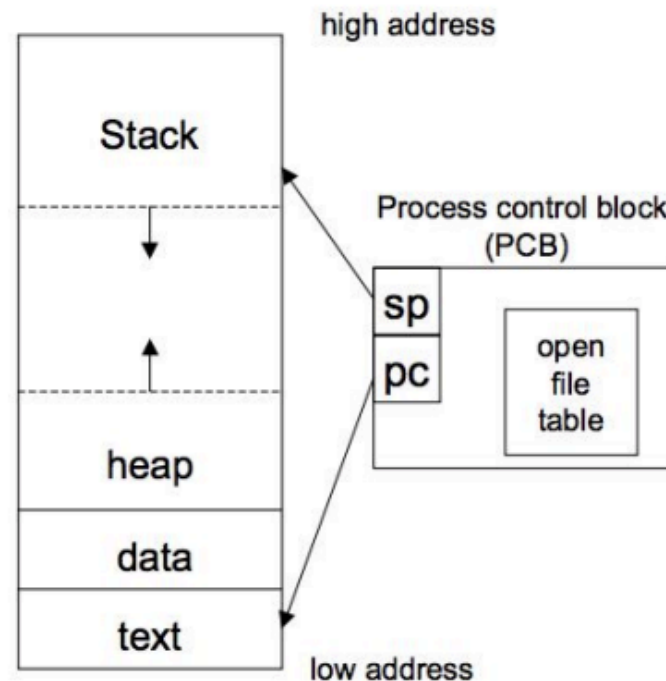
- OS allocates a PCB for it
- OS initializes PCB
- OS puts PCB on the correct queue

As a process computes:

- OS moves its PCB from queue to queue

When a process is terminated:

- PCB may hang around for a while (exit code, etc.)
- eventually, OS deallocates the PCB

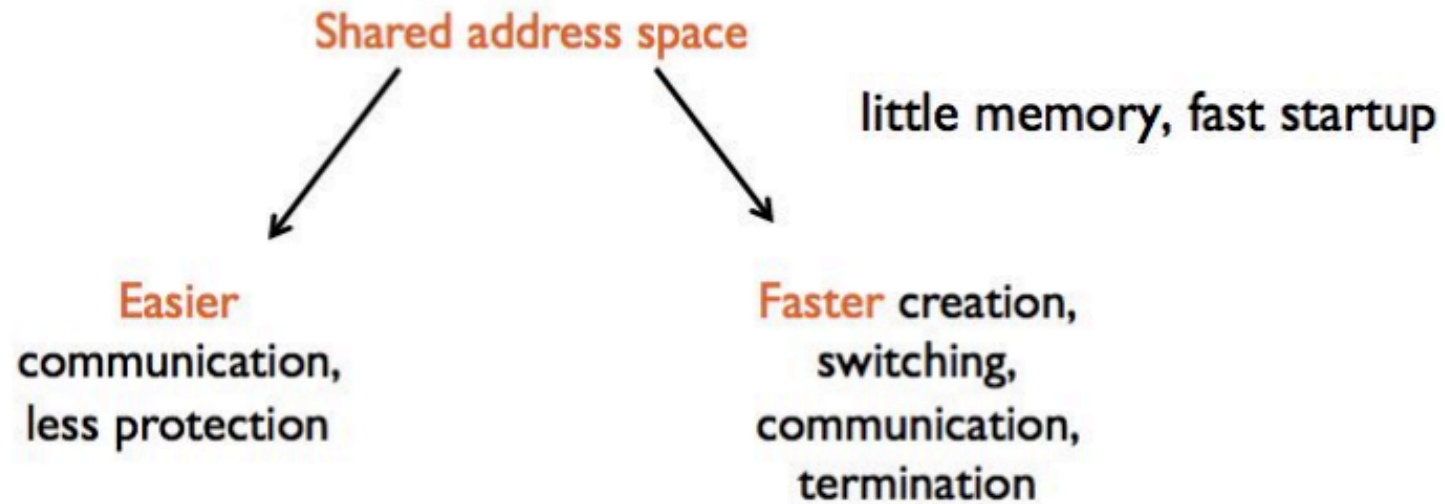


- Each process has its own
 - program counter
 - stack
 - stack pointer
 - address space
- Processes may share
 - open files
 - pipes

<https://courses.cs.washington.edu/courses/cse451/05wi/lectures/4-process.pdf>

Thread

- Light-weight processes



Thread-Properties

- Threads
- Execute in same address space
 - separate execution stack, share access to code and (global) data
- Smaller creation and context-switch time
- Can exploit fine-grain concurrency
- Easier to write programs that use asynchronous I/O or communication

Process vs threads

Processes

- Exploit parallelism successfully
- Separate memory space: good for protection

Threads

- Exploit parallelism successfully
- Shared memory space: good for working together

Thread-continue

- User-level vs kernel-level threads
 - kernel not aware of threads created by user level thread package (e.g. **Pthreads**), language (e.g. Java)
 - user-level threads typically multiplexed on top of kernel level threads in a user-transparent fashion

Implementing Threads in User Space

Threads managed by a threads library


- Kernel is unaware of presence of threads

Advantages:

- No kernel modifications needed to support threads
- Efficient: creation/deletion/switches don't need system calls
- Flexibility in scheduling: library can use different scheduling algorithms, can be application dependent

Disadvantages


- Need to avoid blocking system calls
- Threads compete for one another
- Does not take advantage of multiprocessors [no real parallelism]



A user-level threads package

Kernel Level

- Shared virtual address space
 - Contains running state data
 - Less overhead
 - From the OS's point of view, this is what is scheduled to run on a CPU
-
- No need to create a new address space
 - No need to change address space in context switch
 - Kernel aware
 - Still need to enter kernel to context switch



A threads package managed by the kernel

Thread Creation and Joining

- A thread can be terminated by
 - Returning from the thread function
 - the **main()** function exiting or `exit()` called or sending a **SIGTERM** signal
 - **pthread_exit** - join with a terminated thread
 - **pthread_cancel** - send a cancellation request to a thread
- What is the difference between **exit** and **pthread_exit**?
 - `exit()`: exits the entire process and sets the process's exit value. All threads inside the process are stopped
 - `pthread_exit(void*)`: only stops the calling thread. The pthread library will automatically finish the process if there are no other threads running.
- Passing Arguments to Threads:
<http://www.cs.toronto.edu/~krueger/csc209h/lectures/Week13-threads-4.pdf>