

# CS214 Recitation Sec.7



Oct. 31, 2017

# Topics

1. HW5: Implementing “ps” command in C
2. Differences between threads & processes
3. Thread creation & joining

# HW5 - Implementing "ps"

0. Modify the *ls* we wrote this week except output */proc*

Then open the file '*status*', look for the *uid section* and extract the owner's uid.

Using *pwd.h*, determine the *name of the user* who owns the process and print it out as well.

1. Open *status* alongside/after cmdline, read through and parse the status file looking for 'uid'

# HW5 - Implementing "ps" (Cont.)

2. Then open the file *[schedstat](#)* and read in order: time spent running on CPU (in nanoseconds), time spent waiting on a run queue, # of times context switched

Like default “ps” only prints out info for YOUR procs by default - so can get current uid within your code using *[getuid](#)* and comb through /proc and only print out status and schedstat information for staff whose uid matches

# PS - where is process information

*/proc* - “proc filesystem” is a pseudo-filesystem which provides *an interface to kernel data structures*. It is commonly mounted at */proc*. It includes all the information about processes. The information for each process is stored in a folder name by its PID:

```
-sh-4.2$ ls /proc
```

1	1296	16406	21232	2473	27847	4263	6318	779	9449
10	12972	16442	21247	2475	27859	43	6320	7830	9452
1005	13	16446	21265	24870	27861	4329	6330	784	9453
1011	1307	16453	21267	24990	27883	4331	6331	7841	946
10251	1312	16456	21283	25040	27903	449	6335	785	9462
10353	13327	16465	21295	25088	2793	4551	6350	786	9471
1073	1334	16518	21315	2556	2794	4560	6358	7868	9476

# PS - where is process information

`<pwd.h>` - provides a definition for **struct passwd**, which saves the information such user's login name and program to be used

## NAME

pwd.h - password structure

## SYNOPSIS

```
#include <pwd.h>
```

## DESCRIPTION

The `<pwd.h>` header provides a definition for **struct passwd**, which includes at least the following members:

char	*pw_name	user's <u>login name</u>
uid_t	pw_uid	numerical <u>user ID</u>
gid_t	pw_gid	numerical group ID
char	*pw_dir	initial working directory
char	*pw_shell	<u>program to use as shell</u>

The `gid_t` and `uid_t` types are defined as described in [`<sys/types.h>`](#).

The following are declared as functions and may also be defined as macros. Function prototypes must be provided for use with an ISO C compiler.

```
struct passwd *getpwnam(const char *);
struct passwd *getpwuid(uid_t);
int getpwnam_r(const char *, struct passwd *, char *,
               size_t, struct passwd **);
int getpwuid_r(uid_t, struct passwd *, char *,
               size_t, struct passwd **);
void endpwent(void);
struct passwd *getpwent(void);
void setpwent(void);
```

# PS - where is process information (Cont.)

`/proc/[pid]/stat`

Status information about the process. This is used by `ps(1)`.  
It is defined in the kernel source file `fs/proc/array.c`.

file '`stat`'

The fields, in order, with their proper `scanf(3)` format specifiers, are listed below. Whether or not certain of these fields display valid information is governed by a ptrace access mode `PTRACE_MODE_READ_FSCREDS | PTRACE_MODE_NOAUDIT` check (refer to `ptrace(2)`). If the check denies access, then the field value is displayed as 0. The affected fields are indicated with the marking [PT].

- (1) `pid` %d  
The process ID.
- (2) `comm` %s  
The filename of the executable, in parentheses.  
This is visible whether or not the executable is swapped out.
- (3) `state` %c  
One of the following characters, indicating process state:  
  
R Running  
  
S Sleeping in an interruptible wait

# PS - where is process information (Cont.)

`/proc/[pid]/status`

Provides much of the information in `/proc/[pid]/stat` and `/proc/[pid]/statm` in a format that's easier for humans to parse. Here's an example:

file '*status*'

```
$ cat /proc/$$/status
Name:  bash
Umask: 0022
State: S (sleeping)
Tgid:  17248
Ngid:  0
Pid:   17248
PPid:  17200
TracerPid: 0
Uid:    1000    1000    1000    1000
Gid:    100     100     100     100
FDSize: 256
Groups: 16 33 100
NSTgid: 17248
NSpid:  17248
NSpgid: 17248
NSsid:  17200
VmPeak: 131168 kB
VmSize: 131168 kB
VmLck:  0 kB
VmPin:  0 kB
VmHWM:  13484 kB
VmRSS:  13484 kB
```



# PS - where is process information (Cont.)

file '`comm`'

`/proc/[pid]/comm` (since Linux 2.6.33)

This file exposes the process's `comm` value—that is, the command name associated with the process. Different threads in the same process may have different `comm` values, accessible via `/proc/[pid]/task/[tid]/comm`. A thread may modify its `comm` value, or that of any of other thread in the same thread group (see the discussion of **CLONE\_THREAD** in `clone(2)`), by writing to the file `/proc/self/task/[tid]/comm`. Strings longer than **TASK\_COMM\_LEN** (16) characters are silently truncated.

This file provides a superset of the `prctl(2)` **PR\_SET\_NAME** and **PR\_GET\_NAME** operations, and is employed by `pthread_setname_np(3)` when used to rename threads other than the caller.

# PS - where is process information (Cont.)

file '*schedstat*'

includes time spent running on CPU (in nanoseconds), time spent waiting on a run queue, # of times context switched.

Use the comand */proc/<pid>/schedstat* and get the following output:

```
-sh-4.2$ cat /proc/76/schedstat  
17882 298 2
```

# PS - when time is expressed by jiffies

## The software clock, HZ, and jiffies

The accuracy of various system calls that set timeouts, (e.g., `select(2)`, `sigtimedwait(2)`) and measure CPU time (e.g., `getrusage(2)`) is limited by the resolution of the *software clock*, a clock maintained by the kernel which measures time in *jiffies*. The size of a jiffy is determined by the value of the kernel constant *HZ*.

The value of *HZ* varies across kernel versions and hardware platforms. On i386 the situation is as follows: on kernels up to and including 2.4.x, *HZ* was 100, giving a jiffy value of 0.01 seconds; starting with 2.6.0, *HZ* was raised to 1000, giving a jiffy of 0.001 seconds. Since kernel 2.6.13, the *HZ* value is a kernel configuration parameter and can be 100, 250 (the default) or 1000, yielding a jiffies value of, respectively, 0.01, 0.004, or 0.001 seconds. Since kernel 2.6.20, a further frequency is available: 300, a number that divides evenly for the common video frame rates (PAL, 25 HZ; NTSC, 30 HZ).

The `times(2)` system call is a special case. It reports times with a granularity defined by the kernel constant *USER\_HZ*. User-space applications can determine the value of this constant using `sysconf(_SC_CLK_TCK)`.

# Isdigit function in <ctype.h>

```
#include <ctype.h>
main(){
    char str[] = "123@#FDsP[e?";
    int i;
    for(i = 0; str[i] != 0; i++)
        if(isdigit(str[i]))
            printf("%c is an digit character\n", str[i]);
}
```

**Output:**

1 is an digit character

2 is an digit character

3 is an digit character

# Getuid function in <unistd.h>

## NAME [top](#)

getuid, geteuid - get user identity

## SYNOPSIS [top](#)

```
#include <unistd.h>
#include <sys/types.h>

uid_t getuid(void);
uid_t geteuid(void);
```

## DESCRIPTION [top](#)

**getuid()** returns the real user ID of the calling process.

**geteuid()** returns the effective user ID of the calling process.

# HW5 - A Sample Code

```
#include <ctype.h>
```

```
int check_if_number (char *str)
{
    int i;
    for (i=0; str[i] != '\0'; i++)
    {
        if (!isdigit (str[i]))
        {
            return 0;
        }
    }
    return 1;
}
```

```
#include <pwd.h>
```

```
const char *getUserNamе(int uid)
{
    struct passwd *pw = getpwuid(uid);
    if (pw)
    {
        return pw->pw_name;
    }

    return "";
}
```

## HW5 - A Sample Code (Cont.)

```
#define MAX_BUF 2048
#define INT_SIZE_BUF 8
#define COMM_SIZE_BUF 256
```

```
void pidaux (char *cuser)
```

```
int main (int argc, char *argv[])
{
    char current_user[256];
    strcpy (current_user, getUsername(getuid()));
    printf("UID\t STATE\t PID\t PPID\t NUM_TR\t COMM\t\t WAIT_T\t\t T_SWCH\t\t EXEC_T\n");
    pidaux(current_user);
    long Hertz=sysconf(_SC_CLK_TCK);
    printf("Hertz: %ld\n", Hertz);
    return 0;
}
```

# HW5 - A Sample Code (Cont.)

```
void pidaux (char *cuser)
{
    DIR *thingy;
    FILE *fp;
    struct dirent *entry;

    char path[MAX_BUF], comm_buf[COMM_SIZE_BUF];
    char uid_int_str[INT_SIZE_BUF]={0};
    char *line;
    char *user, *command;
    size_t len=0; // auto allocation
    thingy = opendir ("/proc/");

    if (thingy == NULL)
    {
        perror ("Fail");
        exit(0);
    }
}
```

```
while ((entry = readdir (thingy)) != NULL)
{
    if (check_if_number (entry->d_name))
    {
        strcpy(path, "/proc/");
        strcat(path, entry->d_name);
        strcat(path, "/status");
        fp = fopen(path, "r");

        if(fp!=NULL)
        {
            getline(&line, &len, fp);
            getline(&line, &len, fp);
            getline(&line, &len, fp);
            getline(&line, &len, fp);
            getline(&line, &len, fp);
            getline(&line, &len, fp);
            getline(&line, &len, fp);
            getline(&line, &len, fp);
            sscanf(line, "Uid:    %s ", uid_int_str);
        }
        else
        {
            fprintf(stdout, "FP is NULL\n");
        }
    }
}
```



# HW5 - A Sample Code (Cont.)

```
strcpy(path, "/proc/");  
strcat(path, entry->d_name);  
strcat(path, "/stat");
```

```
fp = fopen(path, "r");  
getline(&line, &len, fp);
```

```
char comm[COMM_SIZE_BUF], state;  
unsigned int flags;  
int pid, ppid, pgrp, session, tty_nr, tpgid;  
unsigned long minflt, cminflt, majflt, cmajflt, utime, stime;  
unsigned long long starttime;  
long cutime, cstime, priority, nice, num_threads, itreavalue;
```

```
sscanf(line, "%d %s %c %d %d %d %d %d %u %lu %lu %lu %lu %lu %lu %ld %ld  
%ld %ld %ld %ld %llu", &pid, comm, &state, &ppid, &pgrp, &session, &tty_nr,  
&tpgid, &flags, &minflt, &cminflt, &majflt, &cmajflt, &utime, &stime, &cutime,  
&cstime, &priority, &nice, &num_threads, &itreavalue, &starttime);
```

# HW5 - A Sample Code (Cont.)

```
strcpy (path, "/proc/");  
strcat (path, entry->d_name);  
strcat (path, "/comm");  
  
fp = fopen (path, "r");  
if (fp != NULL)  
{  
    fscanf(fp, "%s", comm_buf);  
    fclose(fp);  
}
```

```
strcpy(path, "/proc/");  
strcat (path, entry->d_name);  
strcat(path, "/schedstat");  
  
fp=fopen(path, "r");  
char exec_runtime_str[256];  
char waiting_on_queue_str[256];  
char times_switched_str[256];  
if(fp!=NULL)  
{  
    getline(&line,&len,fp);  
    sscanf(line,"%s %s %s", exec_runtime_str, waiting_on_queue_str,  
        times_switched_str);  
}  
  
unsigned long long exec_runtime = atol(exec_runtime_str);  
unsigned long waiting_on_queue = atol(waiting_on_queue_str);  
unsigned long times_switched = atol(times_switched_str);
```

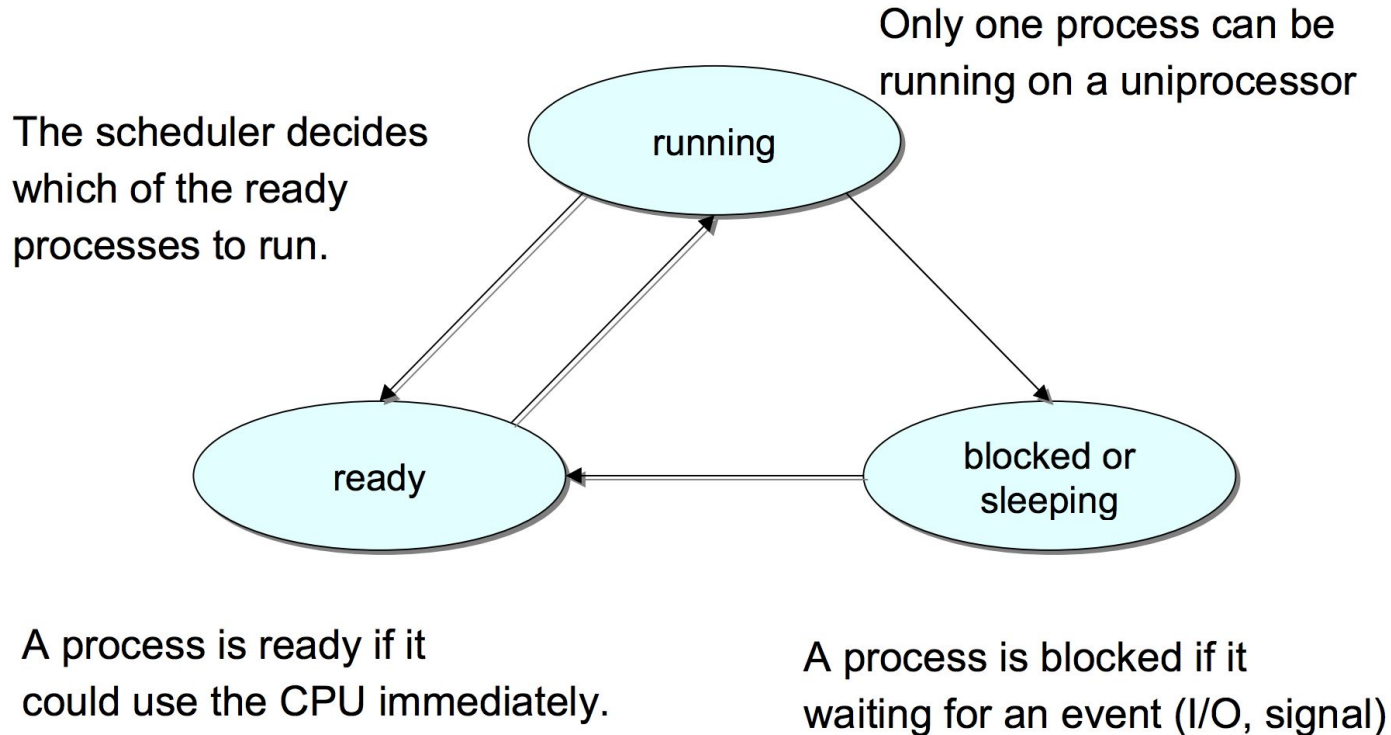
# HW5 - A Sample Code (Cont.)

```
// only print for the current user
if (strcmp(getUserName(atoi(uid_int_str)), cuser) == 0)
{
    if (strlen(getUserName(atoi(uid_int_str))) < 9)
    {
        user = getUserName(atoi(uid_int_str));
    }
    else
    {
        user = uid_int_str;
    }
    printf("%s\t %c\t %d\t %d\t %ld\t %s\t\t %lu\t\t %lu\t\t %llu\n", user, state,
        pid, ppid, num_threads, comm_buf, waiting_on_queue, times_switched, exec_runtime);
}
}
}
closedir (thingy);
}
```

# Recall: Process states

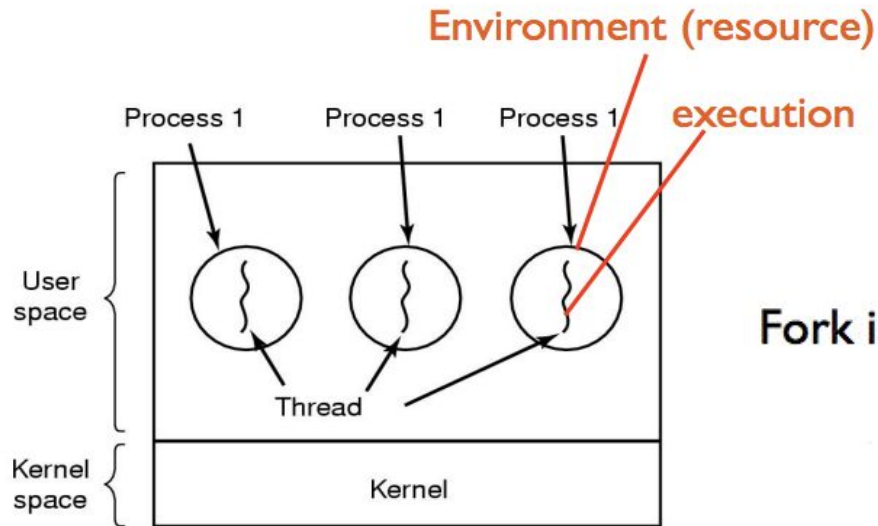
In a single CPU system what is the maximum number of processes that can be in the running state?

# Recall: Process states (main memory)



# Recall: Properties of processes

Processes do not share resources well



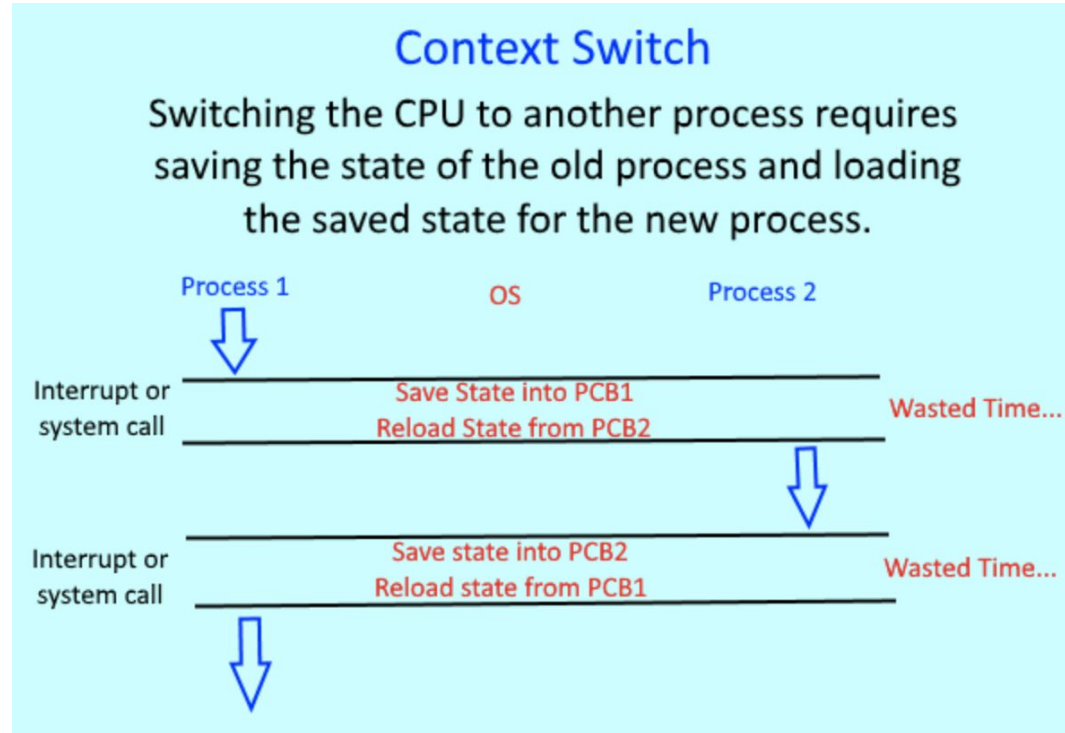
Fork is expensive (time & memory)

Three processes each with  
one thread

# Recall: Properties of processes (Cont.)

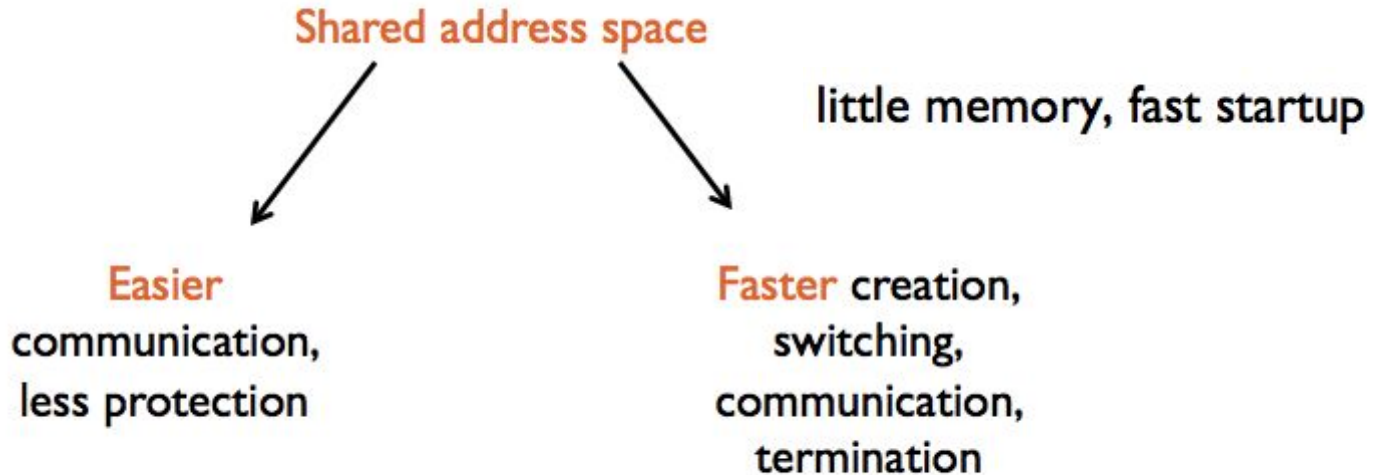
Process context switching cost is high

**Context switch** enables multiple processes to share a single CPU and is an essential feature of a multitasking operating system



# Recall: Properties of threads (Cont.)

Therefore ... Threads: light-weight processes






# Threads

Each process can include many threads.

A **thread** is the smallest entity scheduled for execution on the CPU. *In general, a thread is a component of a process.*

The typical difference is that *threads (of the same process) run in a shared global memory space*, while processes run in separate memory spaces.

Suitable for dealing with *multiple parallel sub-tasks*




# Threads (Cont.)

All threads of a process share:

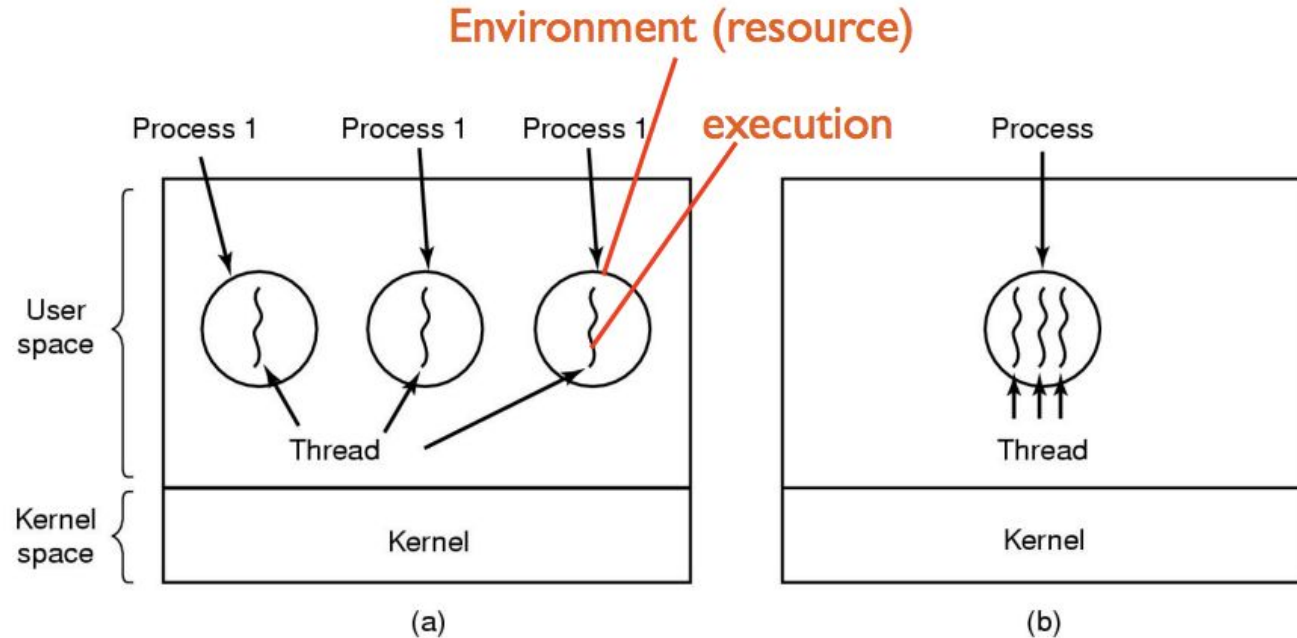
- Process ID
- Memory (program code and global data)
- Working environment (current directory, user ID, etc.)
- Signal handlers

# Threads (Cont.)

- Each has its own function calls & automatic (local) variables
- Need *program counter* and *stack* for each thread
- Each thread has its own *Thread ID* (integer)



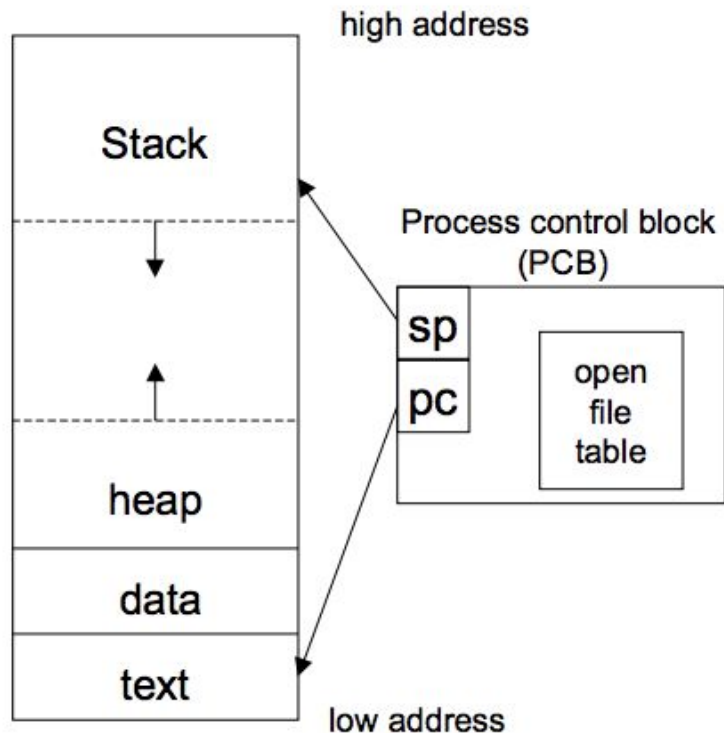
# Differences between threads & processes



Three processes each with  
one thread

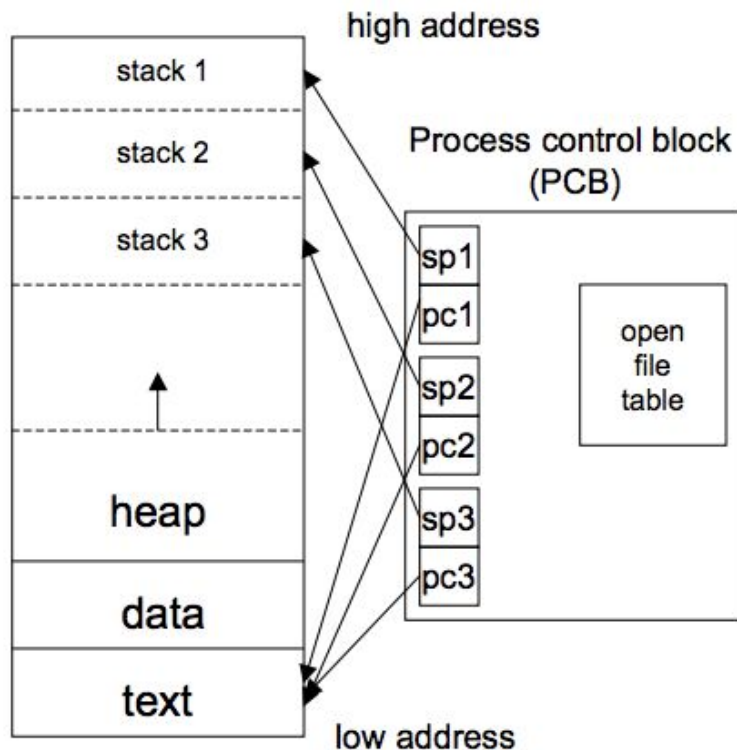
One process with  
three threads

# Memory for processes



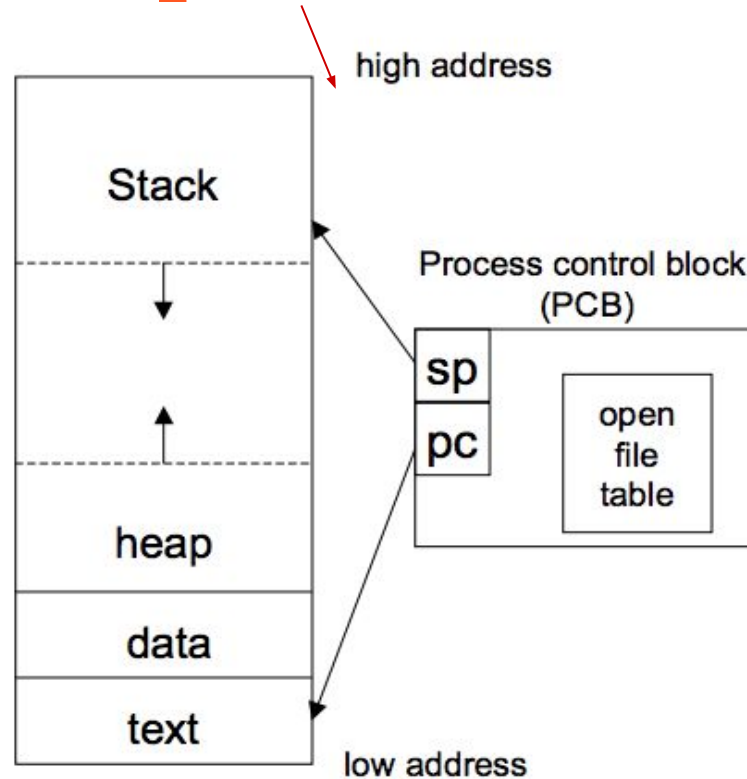
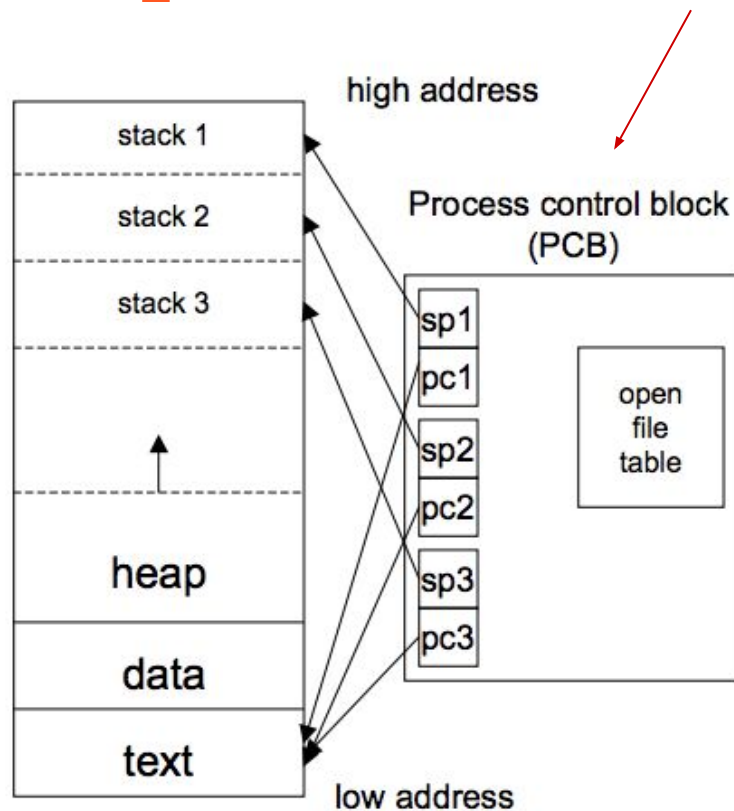
- Each process has its own
  - program counter
  - stack
  - stack pointer
  - address space
- Processes may share
  - open files
  - pipes

# Memory for threads



- Each thread has its own
  - program counter
  - stack
  - stack pointer
- Threads share
  - address space
    - variables
    - code
  - open files

# Comparison of threads & processes



# Creation time for threads & processes

Platform	fork()			pthread_create()		
	real	user	sys	real	user	sys
AMD 2.3 GHz Opteron (16 cpus)	12.5	1.0	12.5	1.2	0.2	1.3
AMD 2.4 GHz Opteron (8 cpus)	17.6	2.2	15.7	1.4	0.3	1.3
IBM 4.0 GHz POWER6 (8 cpus)	9.5	0.6	8.8	1.6	0.1	0.4
IBM 1.9 GHz POWER5 p5-575 (8 cpus)	64.2	30.7	27.6	1.7	0.6	1.1
IBM 1.5 GHz POWER4 (8 cpus)	104.5	48.6	47.2	2.1	1.0	1.5
INTEL 2.4 GHz Xeon (2 cpus)	54.9	1.5	20.8	1.6	0.7	0.9
INTEL 1.4 GHz Itanium2 (4 cpus)	54.5	1.1	22.2	2.0	1.2	0.6

<http://www.llnl.gov/computing/tutorials/pthreads>.

Timings reflect 50,000 process/thread

Creations were performed with the `time` utility, and units are in seconds, no optimization flags.



# Decision: processes or threads?

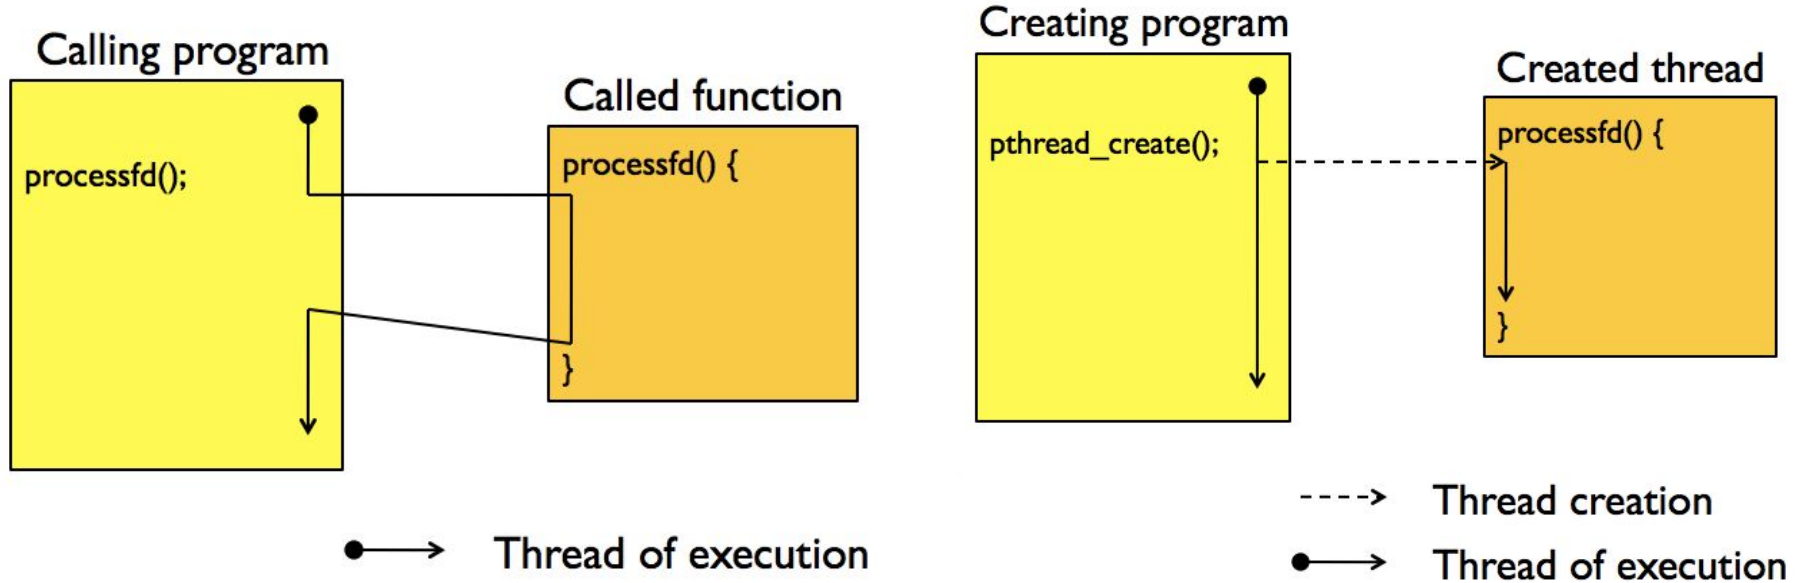
## Processes

- Exploit parallelism successfully
- Separate memory space: good for protection

## Threads

- Exploit parallelism successfully
- Shared memory space: good for working together

# 1 Thread vs Threaded function call



# Pthread

**POSIX threads** (pthreads) is the most commonly used thread package on Unix/Linux

Frequently-used functions in <pthread.h>:

- *pthread\_create* - create a new thread
- *pthread\_join* - join with a terminated thread
- *pthread\_exit* - terminate calling thread
- *pthread\_cancel* - send a cancellation request to a thread

# Thread Creation & Joining

*pthread\_create* - create a new thread

**thread** identifies a thread within a process

**attr** sets attributes such as priority, *initial stack size* – can be specified as *NULL* to get defaults

**start-routine** - the function to call to start the thread

**arg** is the argument to **start-routine**

## SYNOPSIS

[top](#)

```
#include <pthread.h>

int pthread_create(pthread_t *thread, const pthread_attr_t *attr,
                  void *(*start_routine) (void *), void *arg);
```

Compile and link with *-pthread*.

## DESCRIPTION

[top](#)

The **pthread\_create()** function starts a new thread in the calling process. The new thread starts execution by invoking *start\_routine()*; *arg* is passed as the sole argument of *start\_routine()*.

# Thread Creation & Joining

*If I call `pthread_create` twice, how many stacks does my process have?*

Your process will contain three stacks - one for each thread. The first thread is created when the process starts, and you created two more.

# Thread Creation & Joining

*pthread\_join* - join with a terminated thread

**thread** identifies a thread within a process

**retval** returned value by the thread when it terminates

Purpose of pthread\_join():

1. Wait for a thread to finish
2. Clean up thread resources
3. Grabs the return value of the thread

## SYNOPSIS

[top](#)

```
#include <pthread.h>
```

```
int pthread_join(pthread_t thread, void **retval);
```

Compile and link with *-pthread*.

## DESCRIPTION

[top](#)

The **pthread\_join()** function waits for the thread specified by *thread* to terminate. If that thread has already terminated, then **pthread\_join()** returns immediately. The thread specified by *thread* must be joinable.

# Thread Creation & Joining

A thread can be terminated by

- Returning from the thread function
- the *main()* function exiting or *exit()* called or sending a *SIGTERM* signal
- *pthread\_exit* - join with a terminated thread
- *pthread\_cancel* - send a cancellation request to a thread

What is the difference between *exit* and *pthread\_exit*?

- *exit()*: exits the entire process and sets the process's exit value. All threads inside the process are stopped
- *pthread\_exit(void\*)*: only stops the calling thread. The pthread library will automatically finish the process if there are no other threads running.

# Passing Arguments to Threads

```
MyArg *p = (MyArg *)malloc(sizeof(MyArg));
p->fd = fd; /* assumes fd is defined */
strncpy(p->name, "CSC209", 7);
result = pthread_create(&threadID, NULL,
                        myThreadFcn, (void *)p);
void *myThreadFcn(void *p) {
    MyArg *theArg = (MyArg *) p;
    write(theArg->fd, theArg->name, 7);
    close(theArg->fd);
    free(theArg);
    return NULL;
}
```



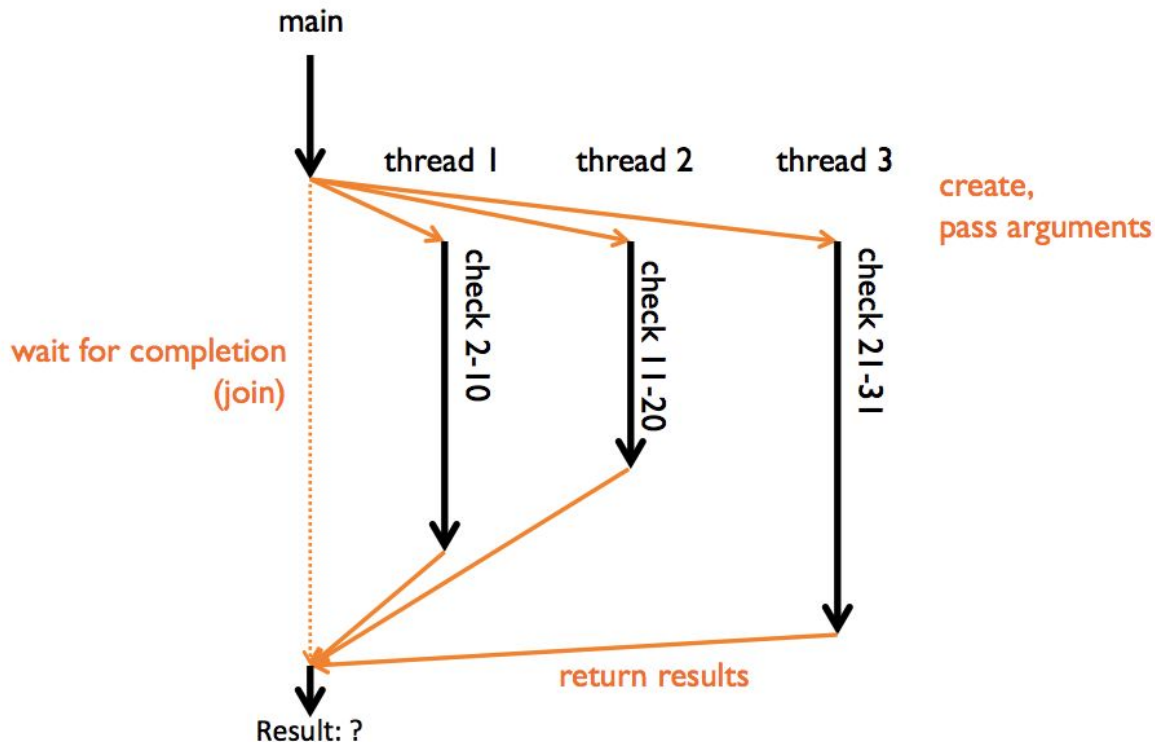
# Thread example

Decide if an integer is prime


- Input: integer
- Output: prime, or composite with factor

Exploit parallelism

- Testing primality can be slow
- My laptop has multiple cores



# Thread example





Happy Halloween