

CS214

Recitation

Sec.7

Nov. 28, 2017

Topics

1. OSI
2. IPv4 vs IPv6
3. TCP vs UDP
4. Sockets

Open Systems Interconnection Model

- A technology standard maintained by the International Standards Organization (ISO)
- An abstract model of how *network protocols* and *equipment* should communicate and work together
- Contains *seven* layers in *two* groups: **Host layers** & **Media layers**
- **Host layers** perform *application-specific functions* such as data formatting, encryption, transmission, and connection management
- **Media layers** provide more primitive *network-specific functions* such as routing, addressing, and flow control

Open Systems Interconnection Model

OSI Model				
	Layer	Data unit	Function ³	Examples
Host layers	7. Application	Data	High-level APIs, including resource sharing, remote file access, directory services and virtual terminals	HTTP, FTP, SMTP
	6. Presentation		Translation of data between a networking service and an application; including character encoding, data compression and encryption/decryption	ASCII, EBCDIC, JPEG
	5. Session		Managing communication sessions, i.e. continuous exchange of information in the form of multiple back-and-forth transmissions between two nodes	RPC, PAP
	4. Transport	Segments	Reliable transmission of data segments between points on a network, including segmentation, acknowledgement and multiplexing	TCP, UDP
Media layers	3. Network	Packet/Datagram	Structuring and managing a multi-node network, including addressing, routing and traffic control	IPv4, IPv6, IPsec, AppleTalk
	2. Data link	Bit/Frame	Reliable transmission of data frames between two nodes connected by a physical layer	PPP, IEEE 802.2, L2TP
	1. Physical	Bit	Transmission and reception of raw bit streams over a physical medium	DSL, USB

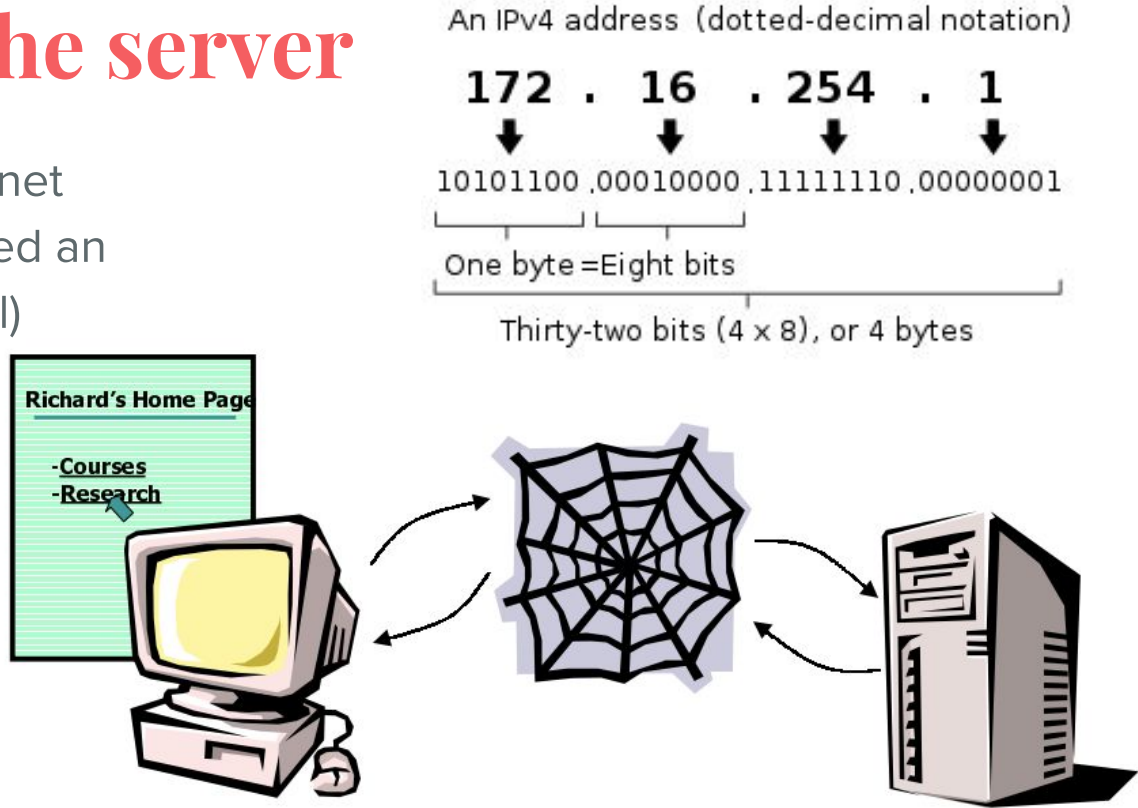
Benefits of OSI

- **Simplifies** the design of network protocols
- Ensure different types of equipment would all be **compatible** if built by different manufacturers
- Makes network designs more **extensible** as new protocols are easier to add to a layered architecture

How do we find the server

- Every computer on the Internet has an Internet address called an **IP** address (Internet Protocol)

- An IP address is four 8-bit numbers separated by dots
- Example: Remote Terminal Server at Rutgers
python.cs.rutgers.edu
128.6.13.233



IPv4 vs IPv6


- **"IPv4"** is version 4 of the Internet Protocol how to send packets of information across a network from one machine to another
- Roughly **95%** of all packets on the Internet today are IPv4 packets
- A significant limitation of IPv4 is that source and destination addresses are limited to **32 bits** (8bits * 4)

IPv4 Header Format

Offsets	Octet	0								1								2								3							
Octet	Bit	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
0	0	Version				IHL				DSCP				ECN				Total Length															
4	32	Identification																Flags				Fragment Offset											
8	64	Time To Live								Protocol								Header Checksum															
12	96	Source IP Address																															
16	128	Destination IP Address																															
20	160	Options (if IHL > 5)																															

IPv4 vs IPv6

- Each IPv4 packet includes a very small header - typically **20 bytes** (more precisely, "octets"), that includes *a source and destination address*
- Conceptually the source and destination addresses can be split into two: a **network number** (the upper bits) and the lower bits represent a **particular host number** on that network.
- A newer packet protocol "**IPv6**" solves many of the limitations of IPv4 (e.g. makes routing tables simpler and **128 bit** addresses) however less than 5% of web traffic is IPv6 based. Example:



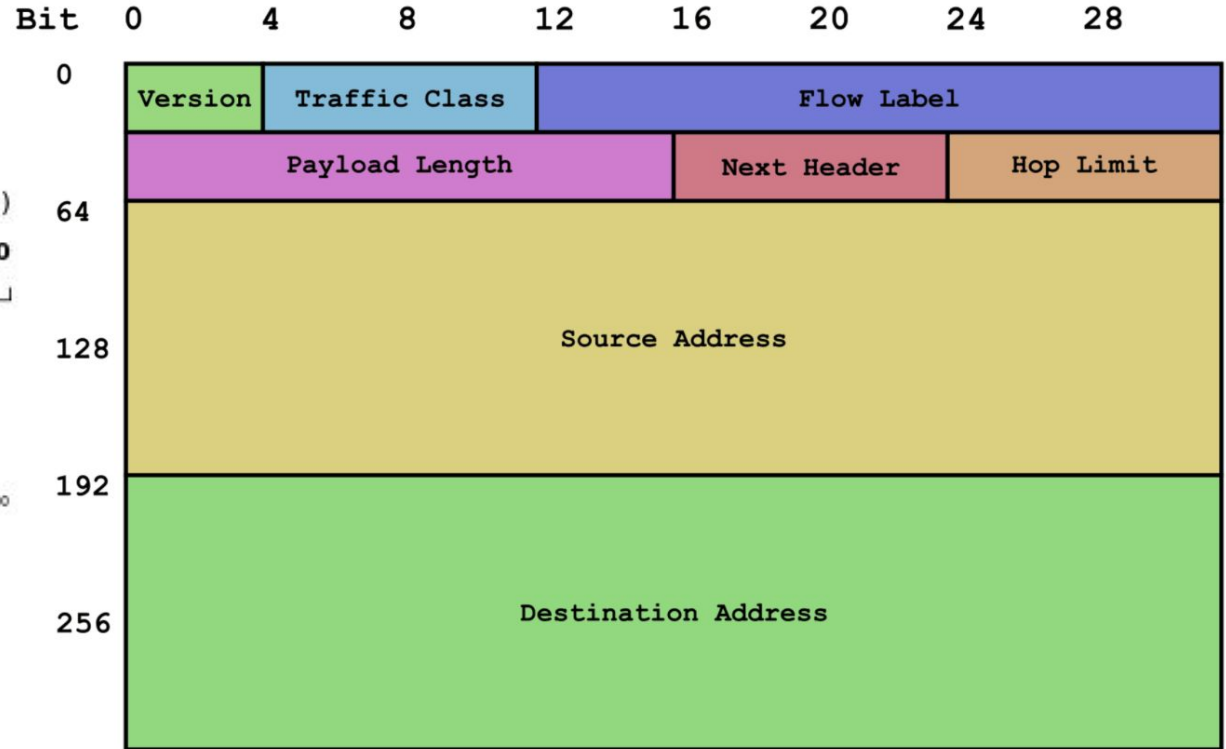
IPv6 packet header

An IPv6 address (in hexadecimal)
2001:0DB8:AC10:FE01:0000:0000:0000:0000

↓ ↓ ↓ ↓

2001:0DB8:AC10:FE01:: Zeroes can be omitted

001000000000000001:0000110110111000:1010110000010000:1111111000000001:
0000000000000000:0000000000000000:0000000000000000:0000000000000000



TCP vs UDP

- UDP is a **connectionless protocol** that is built on top of IPv4 and IPv6
- UDP is very simple to use: 1. Decide the destination address and port and 2. send your data packet
- UDP makes **no guarantee** about whether the packets will arrive
- A typical use case for UDP is when **receiving up to date data is more important** than receiving all of the data

TCP vs UDP

- TCP is a **connection-based protocol** that is built on top of IPv4 and IPv6 (and therefore can be described as "**TCP/IP**" or "TCP over IP")
- TCP creates a **pipe** between two machines and under most conditions, bytes sent from one machine will eventually arrive at the other end without duplication or data loss
- TCP will **automatically manage** resending packets, ignoring duplicate packets, re-arranging out-of-order packets and changing the rate at which packets are sent
- To create a pipe between two machines, TCP use a three-way handshake mechanism which is known as ***SYN, SYN-ACK, and ACK***.

TCP Handshake

Host A **sends** a TCP **SYN**chronize packet to Host B

Host B receives A's **SYN**

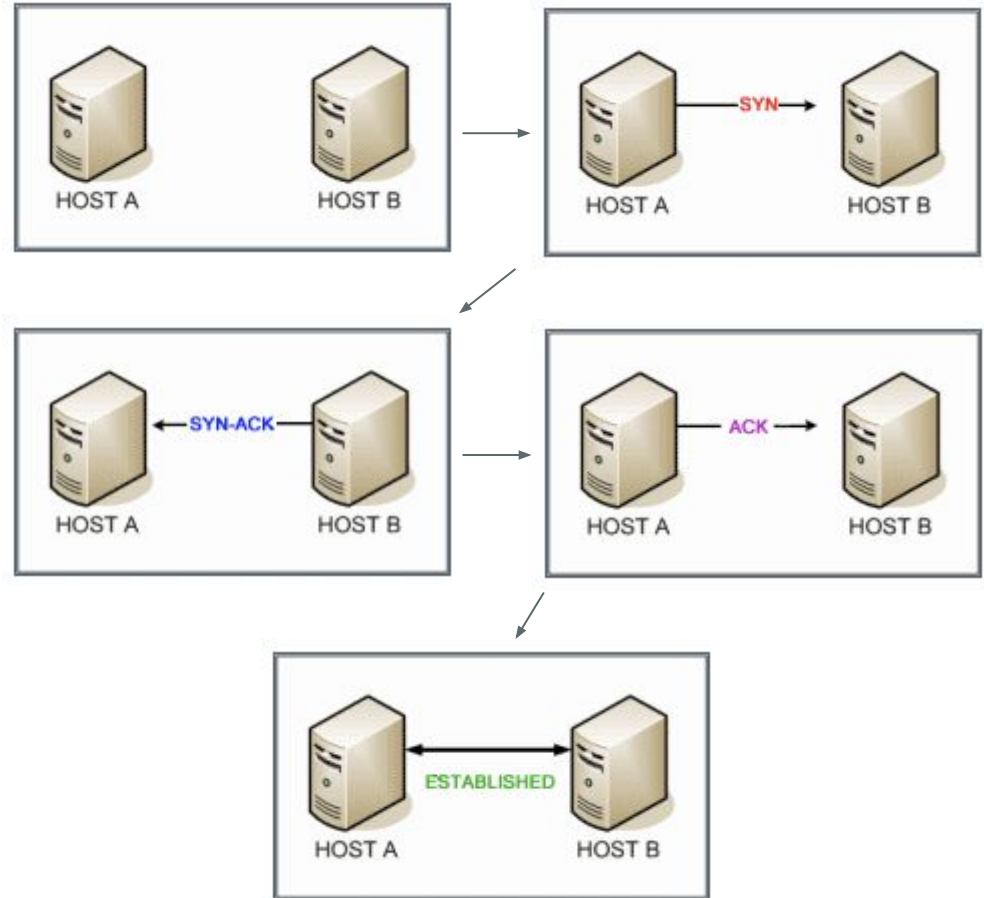
Host B **sends** a **SYN**chronize-**ACK**nowledgement

Host A receives B's **SYN-ACK**


Host A **sends** **ACK**nowledge

Host B receives **ACK**.

TCP socket connection is ESTABLISHED.



Client-Server Model



All communications across a network happen over a **network socket**

- **Client:**
 - **Initiates contact**
 - **Waits for server's response**

- **Server:**
 - **Well-known name**
 - **Waits for contact**
 - **Processes requests, sends replies**

Socket

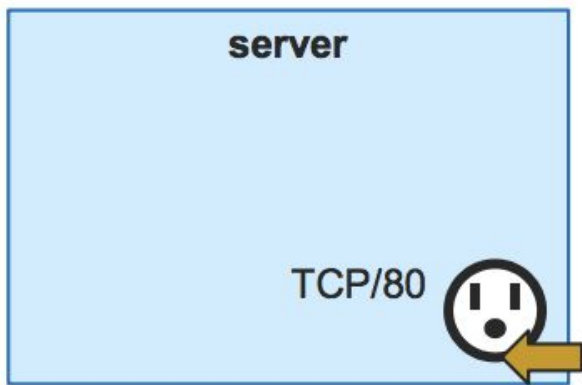
- One form of communication between processes, but it is used between processes on different machines
 - Bi-directional
 - Connection made via a **socket address**
 - **IP address** is the destination of computer
 - **Port number** is the destination of process
- A **socket address** is:
 - IP Address
 - Port Number
 - A socket must also bind to a specific transport-layer protocol.
 - TCP
 - UDP

Socket

- **Port number** is a 16-bit unsigned integer (0 - 65535)
- A *unique resource* shared across the entire system (eg. port 80 can only be utilized by one process)
- Ports below 1024 are *reserved* by operating system
- *Public HTTP servers* always listen for new connections on **port 80**

Initializing a socket

- To listen for an incoming connection, and listen on a specific protocol/port (**Server Socket**)
- To connect to a “server socket” - remote computer (**Client**)



Get socket address with `getaddrinfo`

- The function `getaddrinfo` can convert a human readable domain name (e.g. `www.cs.rutgers.edu`) into an IPv4 and IPv6 address
- A linked-list of `addrinfo` structs will be returned after calling `getaddrinfo`

```
struct addrinfo {  
    int          ai_flags;  
    int          ai_family;  
    int          ai_socktype;  
    int          ai_protocol;  
    socklen_t    ai_addrlen;  
    struct sockaddr *ai_addr;  
    char         *ai_canonname;  
    struct addrinfo *ai_next;  
};
```

Get socket address with getaddrinfo

```
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>

int getaddrinfo(const char *restrict node,
                const char *restrict service,
                const struct addrinfo *restrict hints,
                struct addrinfo **restrict res);
```

■ Parameters

- **node**: host name or IP address to connect to
- **service**: a port number ("80") or the name of a service (found /etc/services: "http")
- **hints**: a filled out struct addrinfo

Using getaddrinfo

output:

```
~/2017F/CS 214/recitation_11_28 » ./addinfo
128.6.68.137
128.6.68.137
```

convert *www.cs.rutgers.edu* into
an IPv4 address with
getaddrinfo

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>

struct addrinfo hints, *infoPtr; // So no need to use memset global variables

int main() {
    hints.ai_family = AF_INET; // AF_INET means IPv4 only addresses

    int result = getaddrinfo("www.cs.rutgers.edu", NULL, &hints, &infoPtr);
    if (result) {
        fprintf(stderr, "getaddrinfo: %s\n", gai_strerror(result));
        exit(1);
    }

    struct addrinfo *p;
    char host[256];

    for(p = infoPtr; p != NULL; p = p->ai_next) {
        getnameinfo(p->ai_addr, p->ai_addrlen, host, sizeof(host), NULL, 0, NI_NUMERICHOST);
        puts(host);
    }

    freeaddrinfo(infoPtr);
    return 0;
}
```

Using getaddrinfo

output:

```
~/2017F/CS 214/recitation_11_28 » ./addinfo  
getaddrinfo: nodename nor servname provided, or not known
```

```
#include <stdio.h>  
#include <stdlib.h>  
#include <sys/types.h>  
#include <sys/socket.h>  
#include <netdb.h>
```

```
struct addrinfo hints, *infoptr; // So no need to use memset global variables
```

convert *www.cs.rutgers.edu* into `int main() {`

an IPv6 address with

`getaddrinfo`

```
hints.ai_family = AF_INET6; // Only want IPv6 (use AF_INET for IPv4)  
hints.ai_socktype = SOCK_STREAM; // Only want stream-based connection
```

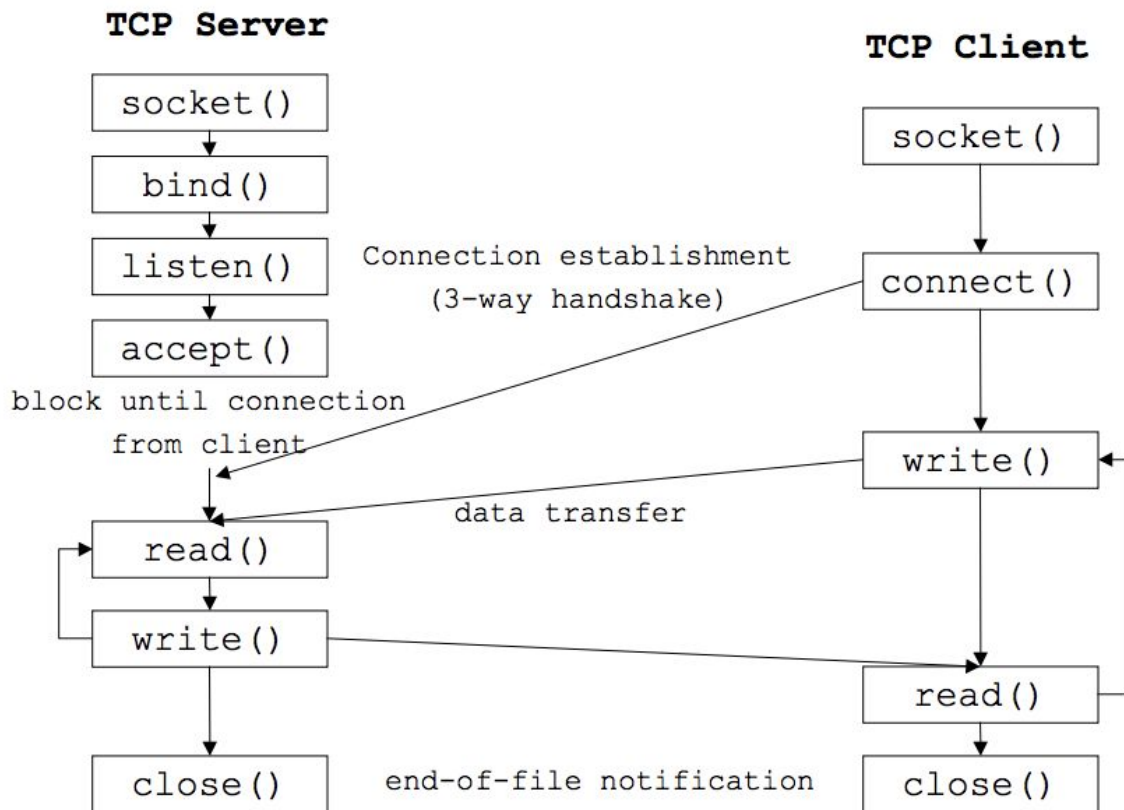
```
int result = getaddrinfo("www.cs.rutgers.edu", NULL, &hints, &infoptr);  
if (result) {  
    fprintf(stderr, "getaddrinfo: %s\n", gai_strerror(result));  
    exit(1);  
}
```

```
struct addrinfo *p;  
char host[256];
```

```
for(p = infoptr; p != NULL; p = p->ai_next) {  
    getnameinfo(p->ai_addr, p->ai_addrlen, host, sizeof(host), NULL, 0, NI_NUMERICHOST);  
    puts(host);  
}
```

```
freeaddrinfo(infoptr);  
return 0;  
}
```

TCP Server-Client Model



Creating a “Server Socket”

- **socket()**: Creates a new socket for a specific protocol (eg: TCP)
- **bind()**: Binds the socket to a specific port (eg: 80)
- **listen()**: Moves the socket into a state of listening for incoming connections.
- **accept()**: Accepts an incoming connection.

Creating a “Client Socket”

- **socket()**: Creates a new socket for a specific protocol (eg: TCP)
- **connect()**: Makes a network connection to a specified IP address and port.

socket()

```
int socket (int family, int type, int protocol);
```

- Create a socket.
 - Returns file descriptor or -1. Also sets **errno** on failure.
 - **family**: address family (namespace)
 - **AF_INET** for IPv4
 - other possibilities: **AF_INET6** (IPv6), **AF_UNIX** or **AF_LOCAL** (Unix socket), **AF_ROUTE** (routing)
 - **type**: style of communication
 - **SOCK_STREAM** for TCP (with **AF_INET**)
 - **SOCK_DGRAM** for UDP (with **AF_INET**)
 - **protocol**: protocol within family
 - typically 0

bind()

```
int bind (int sockfd, struct sockaddr*  
myaddr, int addrlen);
```

- Bind a socket to a local IP address and port number
 - Returns 0 on success, -1 and sets **errno** on failure
 - **sockfd**: socket file descriptor (returned from **socket**)
 - **myaddr**: includes IP address and port number
 - IP address: set by kernel if value passed is **INADDR_ANY**, else set by caller
 - port number: set by kernel if value passed is 0, else set by caller
 - **addrlen**: length of address structure
 - **= sizeof (struct sockaddr_in)**

listen()

```
int listen (int sockfd, int backlog);
```

- Put socket into passive state (wait for connections rather than initiate a connection)
 - Returns 0 on success, -1 and sets **errno** on failure
 - **sockfd**: socket file descriptor (returned from **socket**)
 - **backlog**: bound on length of unaccepted connection queue (connection backlog); kernel will cap, thus better to set high
 - Example:

```
if (listen(sockfd, BACKLOG) == -1) {  
    perror("listen");  
    exit(1);  
}
```

Establishing a Connection

- Include file `<sys/socket.h>`

```
int connect (int sockfd, struct  
             sockaddr* servaddr, int addrlen);
```

- Connect to another socket.

```
int accept (int sockfd, struct sockaddr*  
            cliaddr, int* addrlen);
```

- Accept a new connection. Returns file descriptor or -1.

connect()

```
int connect (int sockfd, struct  
             sockaddr* servaddr, int addrlen);
```

- Connect to another socket.
 - Returns 0 on success, -1 and sets **errno** on failure
 - **sockfd**: socket file descriptor (returned from **socket**)
 - **servaddr**: IP address and port number of server
 - **addrlen**: length of address structure
 - **= sizeof (struct sockaddr_in)**
- Can use with UDP to restrict incoming datagrams and to obtain asynchronous errors

accept()

```
int accept (int sockfd, struct sockaddr* cliaddr,  
            int* addrlen);
```

- Block waiting for a new connection
 - Returns file descriptor or -1 and sets **errno** on failure
 - **sockfd**: socket file descriptor (returned from **socket**)
 - **cliaddr**: IP address and port number of client (returned from call)
 - **addrlen**: length of address structure = pointer to **int** set to **sizeof (struct sockaddr_in)**
- **addrlen** is a **value-result** argument
 - the caller passes the size of the address structure, the kernel returns the size of the client's address (the number of bytes written)

Sending and Receiving Data

```
int send(int sockfd, const void * buf,  
        size_t nbytes, int flags);
```

- Write data to a stream (TCP) or “connected” datagram (UDP) socket.
 - Returns number of bytes written or -1.

```
int recv(int sockfd, void *buf, size_t  
        nbytes, int flags);
```

- Read data from a stream (TCP) or “connected” datagram (UDP) socket.
 - Returns number of bytes read or -1.

send()

```
int send(int sockfd, const void * buf, size_t  
nbytes, int flags);
```

- Send data on a stream (TCP) or “connected” datagram (UDP) socket
 - Returns number of bytes written or -1 and sets **errno** on failure
 - **sockfd**: socket file descriptor (returned from **socket**)
 - **buf**: data buffer
 - **nbytes**: number of bytes to try to write
 - **flags**: control flags
 - MSG_PEEK: get data from the beginning of the receive queue without removing that data from the queue

recv()

```
int recv(int sockfd, void *buf, size_t nbytes,  
int flags);
```

- Read data from a stream (TCP) or “connected” datagram (UDP) socket
 - Returns number of bytes read or -1, sets **errno** on failure
 - Returns 0 if socket closed
 - **sockfd**: socket file descriptor (returned from **socket**)
 - **buf**: data buffer
 - **nbytes**: number of bytes to try to read
 - **flags**: see man page for details; typically use 0

Building a simple TCP Client

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>
#include <unistd.h>

int main(int argc, char **argv)
{
    int s;
    int sock_fd = socket(AF_INET, SOCK_STREAM, 0);

    struct addrinfo hints, *result;
    memset(&hints, 0, sizeof(struct addrinfo));
    hints.ai_family = AF_INET; /* IPv4 only */
    hints.ai_socktype = SOCK_STREAM; /* TCP */

    s = getaddrinfo("www.cs.rutgers.edu", "80", &hints, &result);
    if (s != 0) {
        fprintf(stderr, "getaddrinfo: %s\n", gai_strerror(s));
        exit(1);
    }

    if(connect(sock_fd, result->ai_addr, result->ai_addrlen) == -1){
        perror("connect");
        exit(2);
    }

    char *buffer = "GET / HTTP/1.0\r\n\r\n";
    printf("SENDING: %s", buffer);
    printf("===\n");
```

```
write(sock_fd, buffer, strlen(buffer));

char resp[1000];
int len = read(sock_fd, resp, 999);
resp[len] = '\0';
printf("%s\n", resp);

return 0;
}
```

Building a simple TCP Client

```
~/2017F/CS 214/recitation_11_28 » ./tcp_client
```

```
SENDING: GET / HTTP/1.0
```

```
===
```

```
HTTP/1.1 301 Moved Permanently
```

```
Date: Wed, 29 Nov 2017 17:13:04 GMT
```

```
Server: Apache/2.4.6 (CentOS) OpenSSL/1.0.1e-fips mod_auth_gssapi/1.3.1 mod_auth_  
_kerb/5.4 mod_fcgid/2.3.9 mod_nss/2.4.6 NSS/3.19.1 Basic ECC PHP/5.4.16 SVN/1.7.  
14 mod_wsgi/3.4 Python/2.7.5
```

```
Location: http://www.cs.rutgers.edu/
```

```
Content-Length: 234
```

```
Connection: close
```

```
Content-Type: text/html; charset=iso-8859-1
```

```
<!DOCTYPE HTML PUBLIC "-//IETF//DTD HTML 2.0//EN">
```

```
<html><head>
```

```
<title>301 Moved Permanently</title>
```

```
</head><body>
```

```
<h1>Moved Permanently</h1>
```

```
<p>The document has moved <a href="http://www.cs.rutgers.edu/">here</a>.</p>
```

```
</body></html>
```

Client

Building a simple TCP Server

```
#include <string.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>
#include <unistd.h>
#include <arpa/inet.h>

int main(int argc, char **argv)
{
    int s;
    int sock_fd = socket(AF_INET, SOCK_STREAM, 0);

    struct addrinfo hints, *result;
    memset(&hints, 0, sizeof(struct addrinfo));
    hints.ai_family = AF_INET;
    hints.ai_socktype = SOCK_STREAM;
    hints.ai_flags = AI_PASSIVE;
```

```
s = getaddrinfo(NULL, "1234", &hints, &result);
if (s != 0) {
    fprintf(stderr, "getaddrinfo: %s\n", gai_strerror(s));
    exit(1);
}

if (bind(sock_fd, result->ai_addr, result->ai_addrlen) != 0) {
    perror("bind()");
    exit(1);
}
```

```
if (listen(sock_fd, 10) != 0) {
    perror("listen()");
    exit(1);
}

struct sockaddr_in *result_addr = (struct sockaddr_in *) result->ai_addr;
printf("Listening on file descriptor %d, port %d\n", sock_fd, ntohs(result_addr->sin_port));

printf("Waiting for connection...\n");
int client_fd = accept(sock_fd, NULL, NULL);
printf("Connection made: client_fd=%d\n", client_fd);

char buffer[1000];
int len = read(client_fd, buffer, sizeof(buffer) - 1);
buffer[len] = '\0';

printf("Read %d chars\n", len);
printf("===\n");
printf("%s\n", buffer);

return 0;
}
```

Building a simple TCP Server

```
~/2017F/CS 214/recitation_11_28 » ./tcp_server  
Listening on file descriptor 3, port 1234  
Waiting for connection...
```